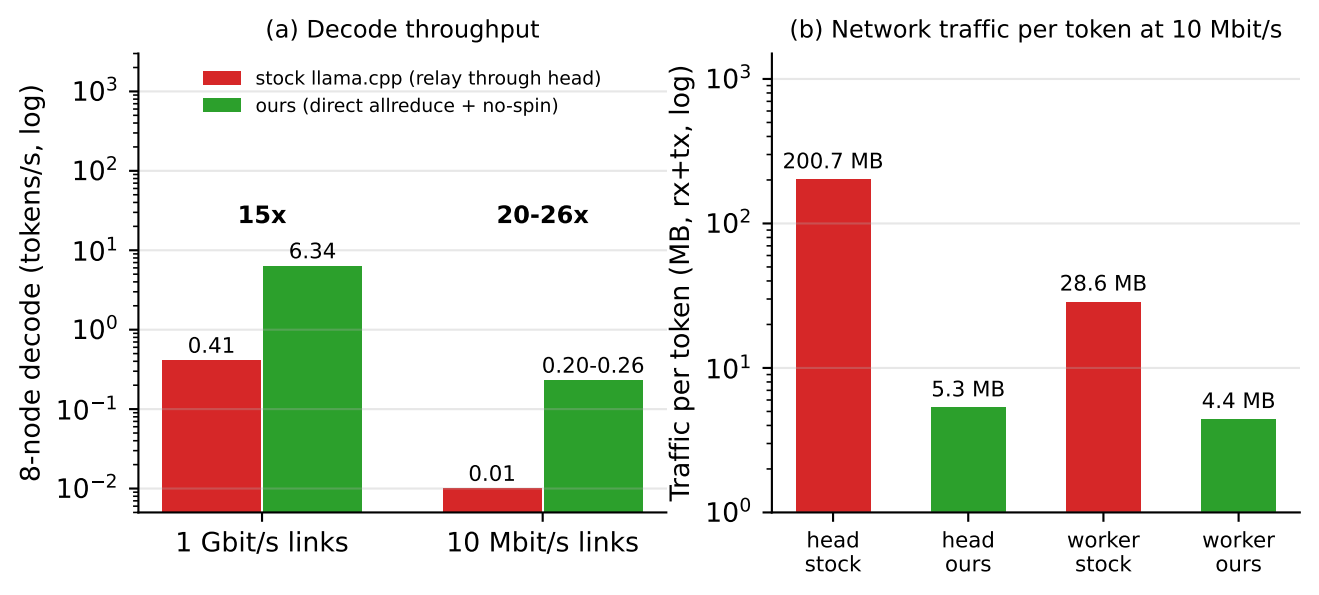
(a) Decode throughput
stock llama.cpp (relay through head)
ours (direct allreduce + no-spin)
8-node decode (tokens/s, log)
10^3
10^2
10^1
10^0
10^{-1}
10^{-2}
15x
20-26x
6.34
0.41
0.20-0.26
0.01
1 Gbit/s links
10 Mbit/s links
(b) Network traffic per token at 10 Mbit/s
Traffic per token (MB, rx+tx, log)
10^3
10^2
10^1
10^0
200.7 MB
5.3 MB
28.6 MB
4.4 MB
head stock
head ours
worker stock
worker ours