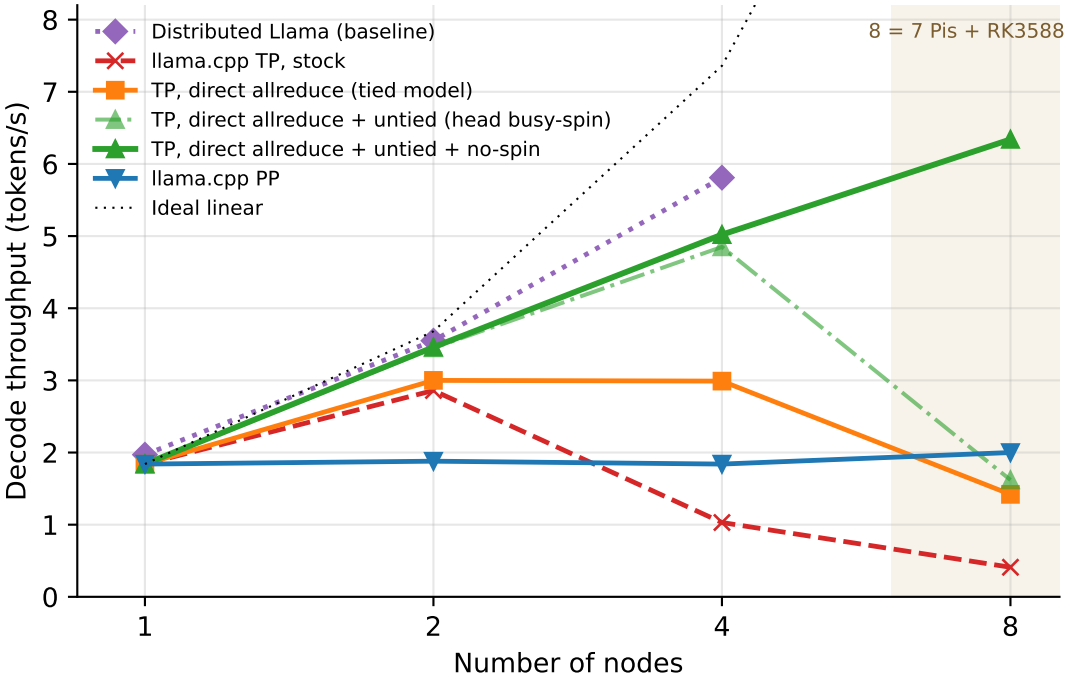
Distributed Llama (baseline)
llama.cpp TP, stock
TP, direct allreduce (tied model)
TP, direct allreduce + untied (head busy-spin)
TP, direct allreduce + untied + no-spin
llama.cpp PP
Ideal linear
8 = 7 Pis + RK3588
Decode throughput (tokens/s)
Number of nodes
0
1
2
3
4
5
6
7
8
1
2
4
8