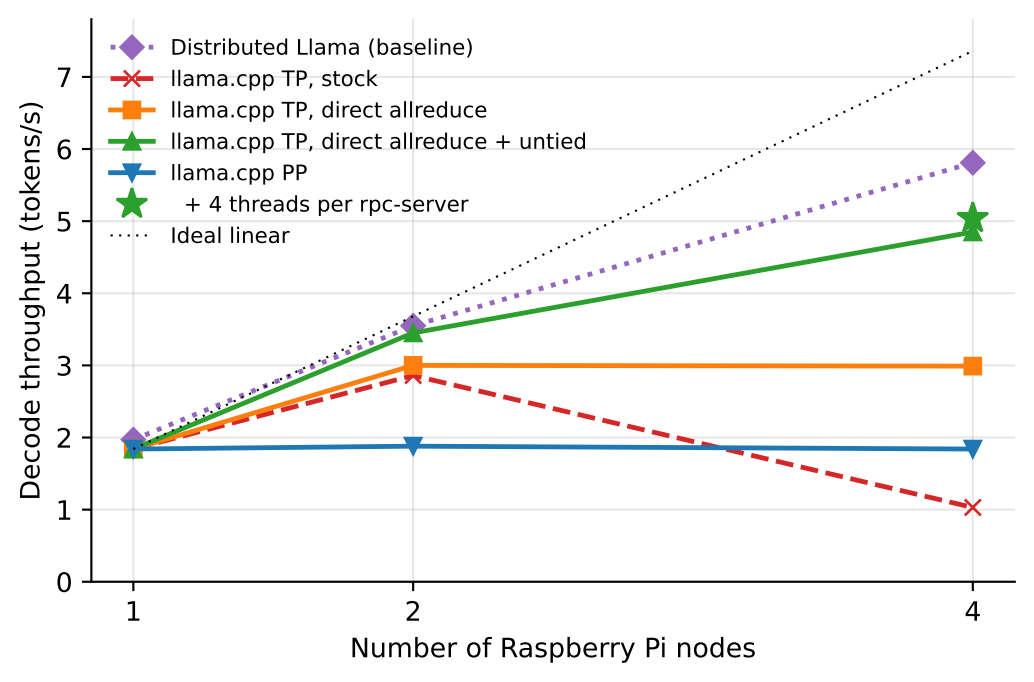
Distributed Llama (baseline)
llama.cpp TP, stock
llama.cpp TP, direct allreduce
llama.cpp TP, direct allreduce + untied
llama.cpp PP
+ 4 threads per rpc-server
Ideal linear
Decode throughput (tokens/s)
0
1
2
3
4
5
6
7
1
2
4
Number of Raspberry Pi nodes