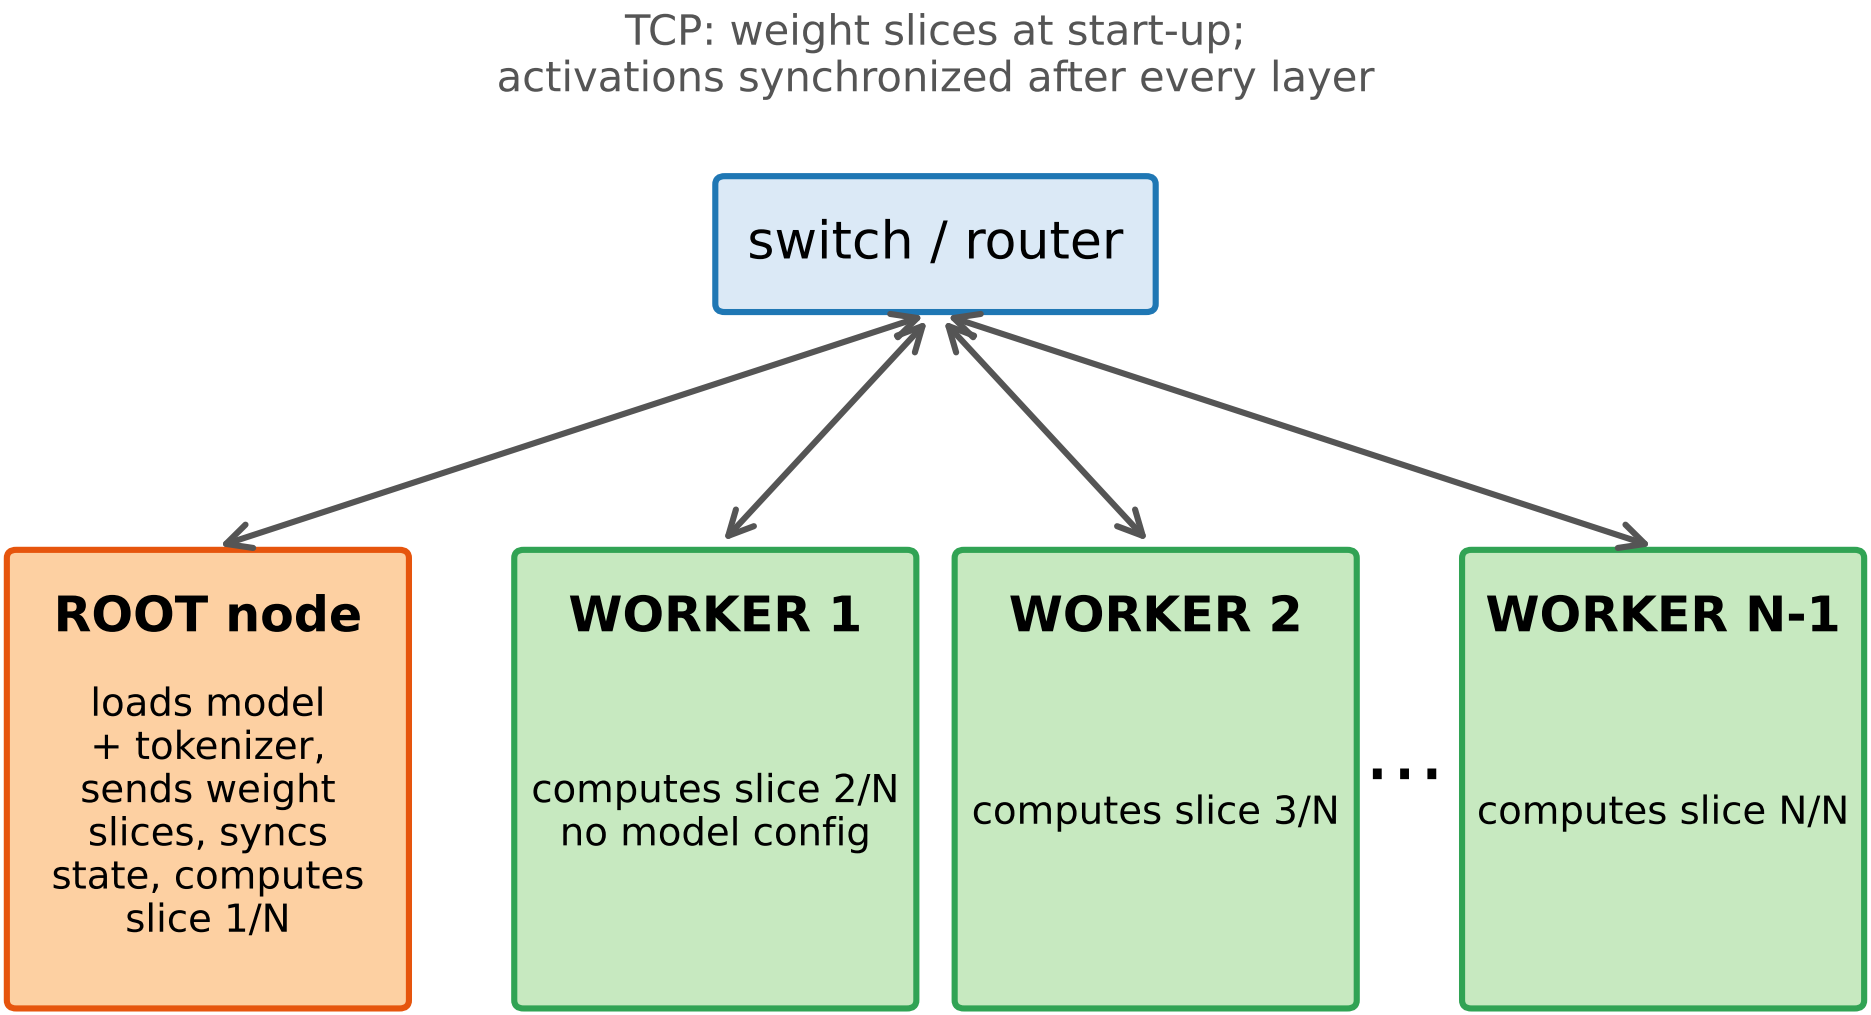
TCP: weight slices at start-up;
activations synchronized after every layer
switch / router
ROOT node
loads model
+ tokenizer,
sends weight
slices, syncs
state, computes
slice 1/N
WORKER 1
computes slice 2/N
no model config
WORKER 2
computes slice 3/N
...
WORKER N-1
computes slice N/N