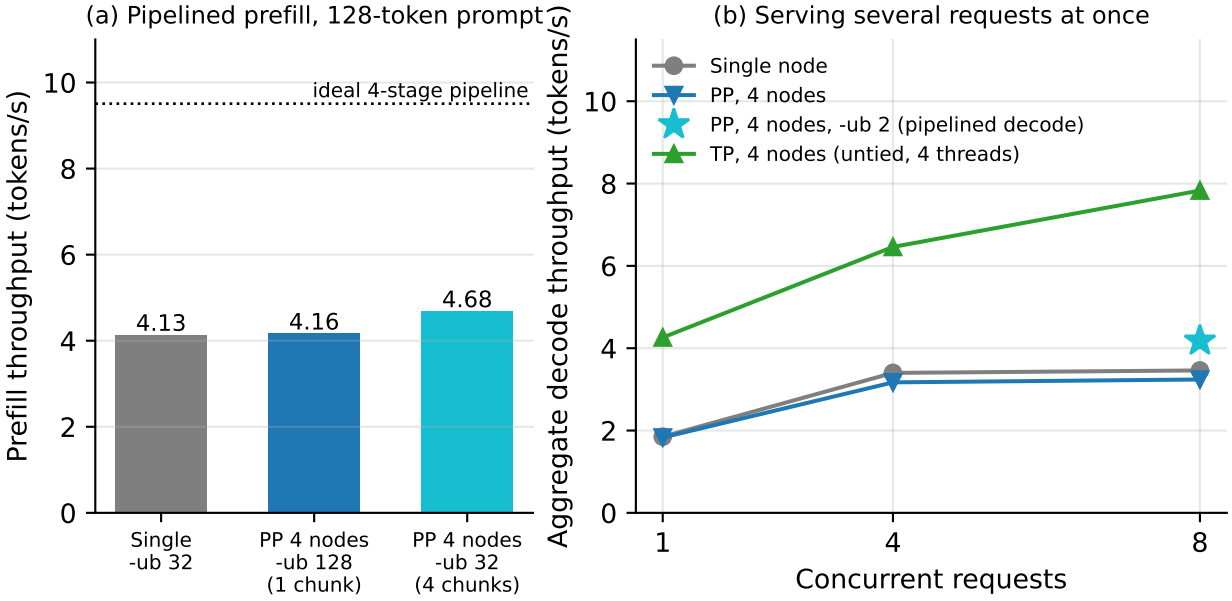
(a) Pipelined prefill, 128-token prompt
Prefill throughput (tokens/s)
ideal 4-stage pipeline
4.13
4.16
4.68
Single
-ub 32
PP 4 nodes
-ub 128
(1 chunk)
PP 4 nodes
-ub 32
(4 chunks)
(b) Serving several requests at once
Aggregate decode throughput (tokens/s)
Single node
PP, 4 nodes
PP, 4 nodes, -ub 2 (pipelined decode)
TP, 4 nodes (untied, 4 threads)
Concurrent requests