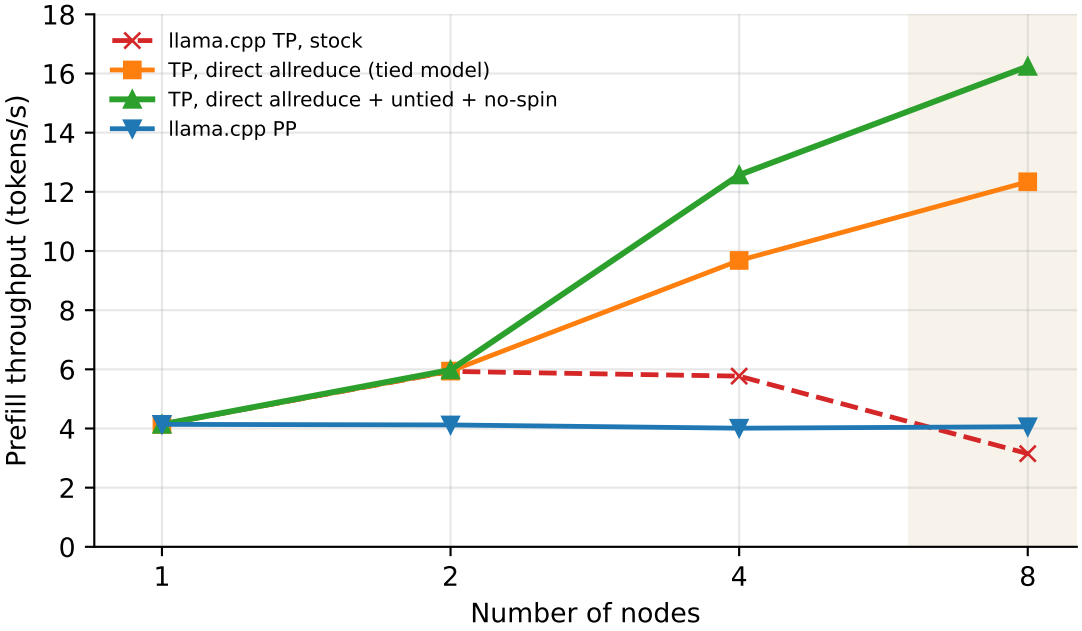
llama.cpp TP, stock
TP, direct allreduce (tied model)
TP, direct allreduce + untied + no-spin
llama.cpp PP
Prefill throughput (tokens/s)
Number of nodes
0
2
4
6
8
10
12
14
16
18
1
2
4
8