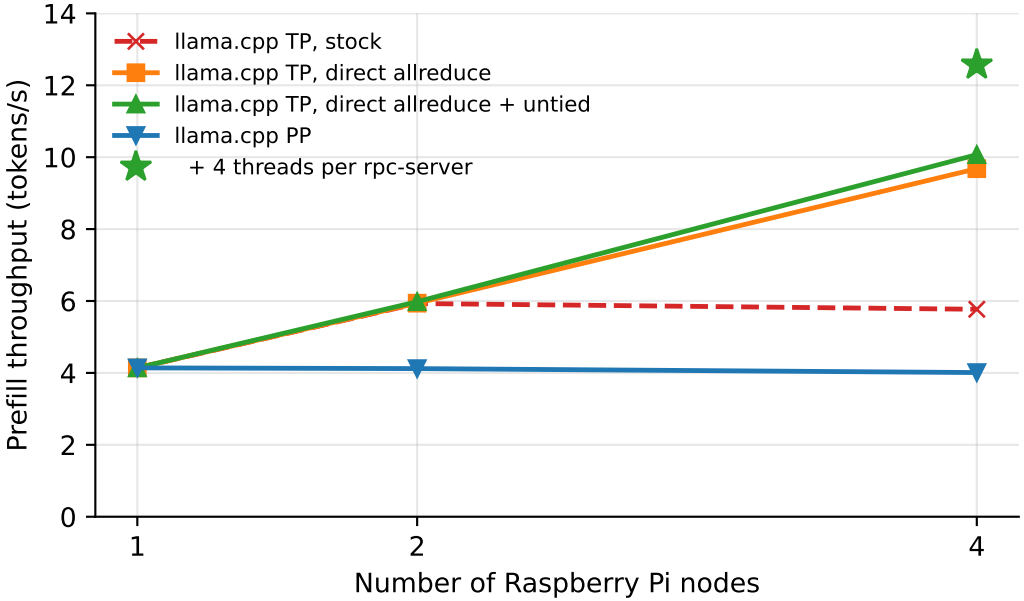
llama.cpp TP, stock
llama.cpp TP, direct allreduce
llama.cpp TP, direct allreduce + untied
llama.cpp PP
+ 4 threads per rpc-server
Prefill throughput (tokens/s)
0
2
4
6
8
10
12
14
1
2
4
Number of Raspberry Pi nodes