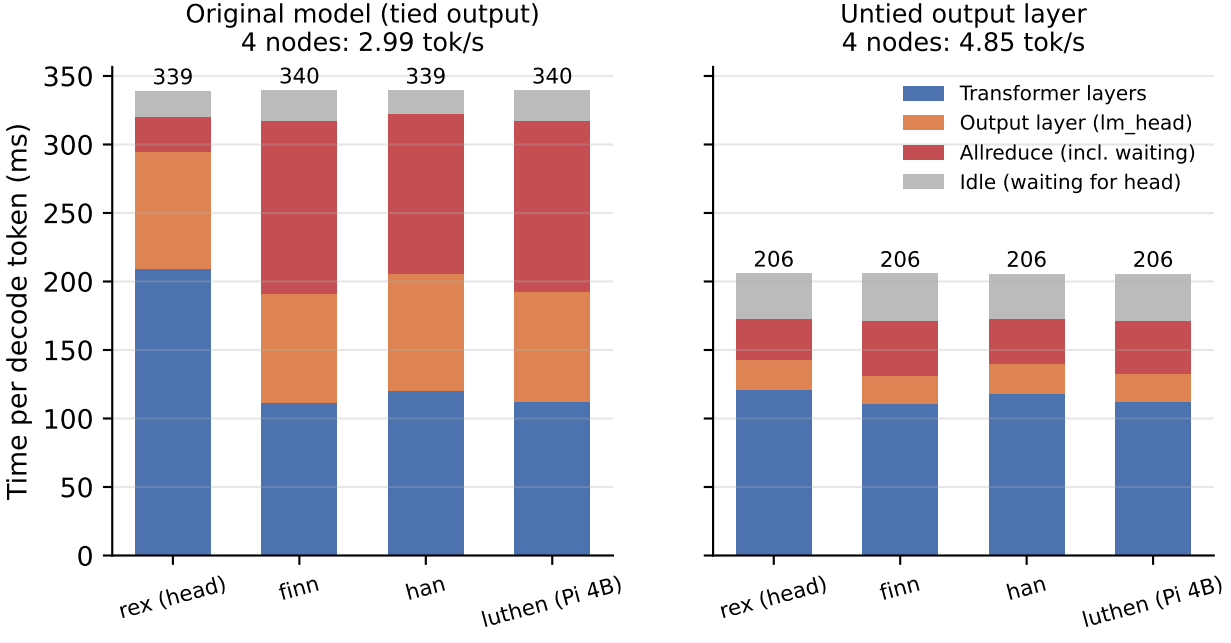
Original model (tied output)
4 nodes: 2.99 tok/s
Time per decode token (ms)
0
50
100
150
200
250
300
350
339
340
339
340
rex (head)
finn
han
luthen (Pi 4B)
Untied output layer
4 nodes: 4.85 tok/s
Transformer layers
Output layer (lm_head)
Allreduce (incl. waiting)
Idle (waiting for head)
206
206
206
206
rex (head)
finn
han
luthen (Pi 4B)