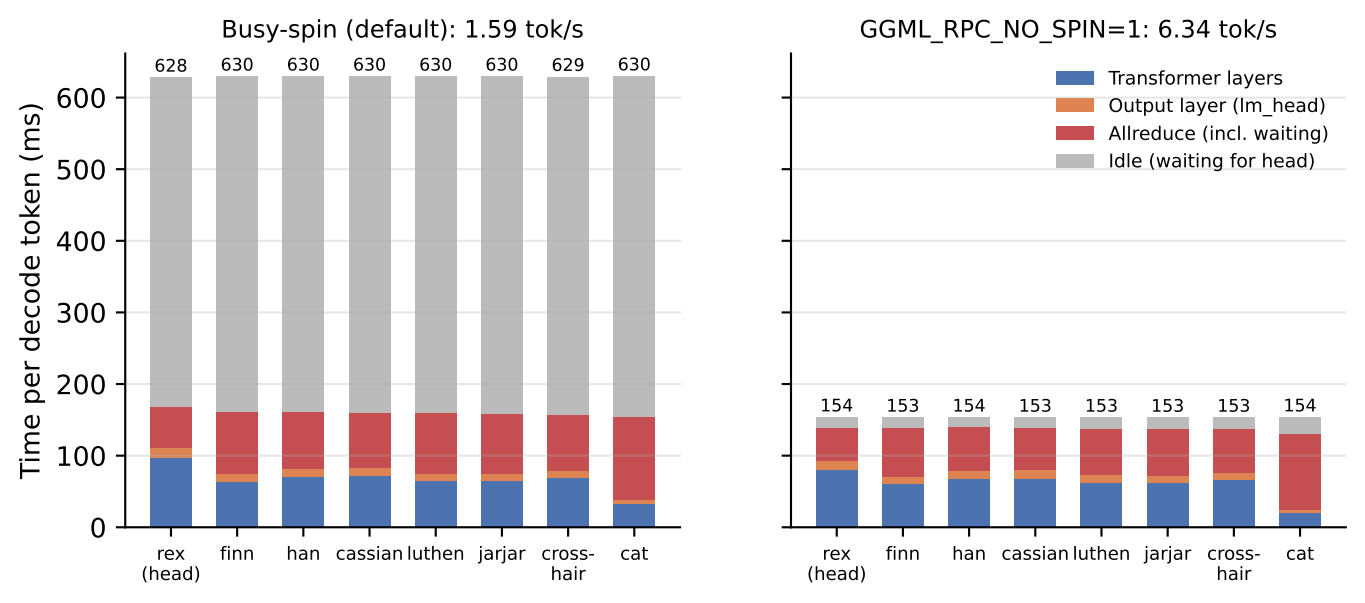
Busy-spin (default): 1.59 tok/s
GGML_RPC_NO_SPIN=1: 6.34 tok/s
Time per decode token (ms)
0
100
200
300
400
500
600
628
630
630
630
630
630
629
630
154
153
154
153
153
153
153
154
Transformer layers
Output layer (lm_head)
Allreduce (incl. waiting)
Idle (waiting for head)
rex (head)
finn
han
cassian
luthen
jarjar
cross-hair
cat